


Blind null model (τ_{blind})

uniform tactic pick over candidates




τ_{blind} = expected attempted edges
to first terminal proof

$\tau_{blind} \approx 120$ attempts

Wonton agent (τ_{agent})

policy + value priors focus search



τ_{agent} = realized attempted edges
on run trace before solve

$\tau_{agent} = 9$ attempts